BNN_FCC Accelerator Design Contest

Pawin Ruangkanit
Course or Lab
Gainesville, Florida, USA
ruangkanitpawin@ufl.edu

1 Summary

1.1 Results

Table 1 summarizes the explored design points. The first three columns specify the architectural parameters used in each run: input bus width, parallel input consumption, and per-layer neuron partitioning. Several design points exceed the maximum operating frequency typically associated with BRAM on the target device. However, these results were achieved using LUTRAM rather than BRAM, indicating that the datapath itself is capable of higher clock frequencies when not limited by BRAM timing. Furthermore, all results were obtained from an out-of-context implementation, which excludes board-level I/O and other full-system constraints. Consequently, these frequencies should be interpreted as optimistic upper-bound estimates rather than fully realizable system-level results. Following the FINN-style reporting convention, FPS is computed from the steady-state throughput rather than the end-to-end latency:

$$\text{FPS} = \frac{F_{\max}}{\text{cycles/output}}$$

where $F_{\max}$ is in Hz. The latency column is still reported separately as the average cycles per image. Among the explored points, the best measured implementation was the 256-bit / 128-input / 64, 64, 10 configuration, which reached 15.56M FPS. The FINN paper reports a maximum of 12.3 million image classifications per second on MNIST, so this best measured point is 1.27× higher, or about 26.5% above that published result. This comparison should still be interpreted carefully, since the values reported here come from an out-of-context LUTRAM-based implementation rather than a full system build.

Table 1: Current design-space results summary.

Input Bus Width	Parallel Inputs	Parallel Neurons	FMAX	Cycles/Output	Latency	FPS
64	32	64,32,10	722	104.8	290	6.89M
128	64	64,32,10	635	56.1	195.2	11.32M
256	128	64,64,10	501	32.2	130.2	15.56M
512	128	256,64,10	—	17.8	95.2	—
1024	128	256,64,10	—	10.6	78.4	—

Table 2: Resource utilization for the top three measured design points.

Input Bus Width	Parallel Inputs	Parallel Neurons	LUT	Registers	LUTRAM	BRAM
64	32	64,32,10	27723 (1.60%)	46757 (1.35%)	8625 (0.50%)	0 (0.00%)
128	64	64,32,10	42518 (2.46%)	60436 (1.75%)	8281 (0.48%)	0 (0.00%)
256	128	64,64,10	67663 (3.92%)	107002 (3.10%)	9037 (0.52%)	0 (0.00%)

Important Judging Notes. Use -p 1.25 in OpenFLEX when building the design. I found this to be most consistent as Vivado gives up early otherwise. Also, I was not able to get the openFLEX timing to work, as I am on a Windows device and the commands did not work, I tried to patch the tool to make it work on windows, but then Modelsim would hang forever. To verify results please use make sim-original-tb.

1.2 Implementation and Tradeoffs

The implementation target was to maximize throughput first and then optimize for clock frequency. Another design priority was parameterization. I wanted every component to be parameterized to be fully scalable to any power-of-two amount of parallel neurons, parallel inputs, and input width. This allowed me to gather these results for various configurations as I scaled up the input bus width.

Throughput is increased by replicating neuron processors and the corresponding weight and threshold RAM banks, allowing more neurons to be processed in parallel each cycle. However, that same replication creates difficult timing paths because the RAM-side control and data signals become large fanout nets with significant routing cost. To address this, multiple pipeline stages were inserted on the interfaces that communicate with the weight and threshold RAM units, breaking these long paths into shorter stages and improving timing closure.

A similar approach is used in the compute reductions. The XNOR accumulation and argmax logic are implemented as recursive balanced trees, which makes it natural to pipeline the computation between tree levels. This reduces the logic depth of each stage and allows the design to scale to wider buses and larger parallel neuron counts without collapsing FMAX. The design is also heavily parameterized, which made it possible to evaluate the architecture across a range of bus widths, parallel-input counts, and neuron configurations while preserving the same overall structure.

The design also supports functionality beyond what was tested for the project. Notably, the design supports partial configuration and input buses as specified by TKEEP, it also supports complete reconfiguration of weights and thresholds in between image streams, as well as customizable hidden layers and topologies.

1.3 Verification/Testbench

The base testbench was extended into a UVM flow to verify functionality beyond the nominal contest path and to provide more complete coverage of the design. Some of this behavior includes: partial TKEEP configuration and input streams, interleaved and reordered configuration traffic, weights-only and thresholds-only reconfiguration, partial-layer updates, deeper hidden-layer topologies, extreme threshold/weight/pixel values, long inter-transaction gaps, output backpressure, and reset/reconfigure behavior.

- Functional coverage was collected with `make coverage-sweep`.
- A text coverage report was generated with `make coverage-sweep-report`.
- The achieved-coverage report is written to `coverage/bnn_fcc_coverage_sweep_test_coverage.txt`.
- Supports partial TKEEP configuration and input streams.
- Supports reconfiguring weights and thresholds after a few images.
- Supports interleaved configuration streams, including thresholds-first and out-of-order layer updates.
- Supports deeper hidden-layer stacks; see `make sim-ten-hidden`.
- Tested long delay gaps and extreme weight, threshold, and pixel values.
- Scales to higher neuron counts when the input bus width increases.

The coverage sweep achieved 93.18% total covergroup coverage, covering 392 of 445 bins across 26 covergroup types, with 0 assertion failures in the reported run.

Table 3: Coverage summary from the directed UVM sweep.

Area	Measured Groups	Status	Main Open Bins
Config stream	toggle 100%, interface 100%, handshake 100%, packet 75.00%, header 80.00%, order 86.66%	Strong	Short or single config packets, plus a few header and ordering corner bins.
Image input	toggle 100%, interface 100%, handshake 100%, spacing 100%, image 62.50%	Strong	Only generic short-packet image-shape bins remain; the contest image path is covered.
Output stream	toggle 100%, interface 100%, backpressure 100%, pattern 100%, output 95.00%	Strong	All 10 classes were observed; remaining misses are a few first-class transition bins.
Config data diversity	weight density 100%, threshold 86.66%	Strong	Main miss is the threshold unknown ratio bin.
Computational stimulus	output class 100%, class transition cross 70.00%	Strong	Every class was produced; misses are mostly “first output was class X” bins.
Reconfiguration	reconfig 83.33%	Good	Full, weights-only, thresholds-only, and partial reconfiguration were hit; the open bin is reconfiguring all layers at once.
Reset / recovery	reset 53.75%, post-reset 100%	Partial	Reset timing was mostly exercised during configuration, not idle, image, output, or TLAST boundaries.

2 Top-Level Organization

This report is organized in the same order that the design is used in practice. The top level accepts a configuration stream that loads weights and thresholds, an input stream that carries activation data, and an output path that returns the classification result. The main body then follows two phases of operation: the configuration phase and the compute phase.

At the top level, `rtl/bnn_fcc.sv` connects three major subsystems: `config_manager`, `data_in_manager`, and `bnn`. The first prepares and distributes model parameters, the second reshapes incoming input data into the binary words expected by the network, and the third contains the replicated layer datapaths that perform inference.

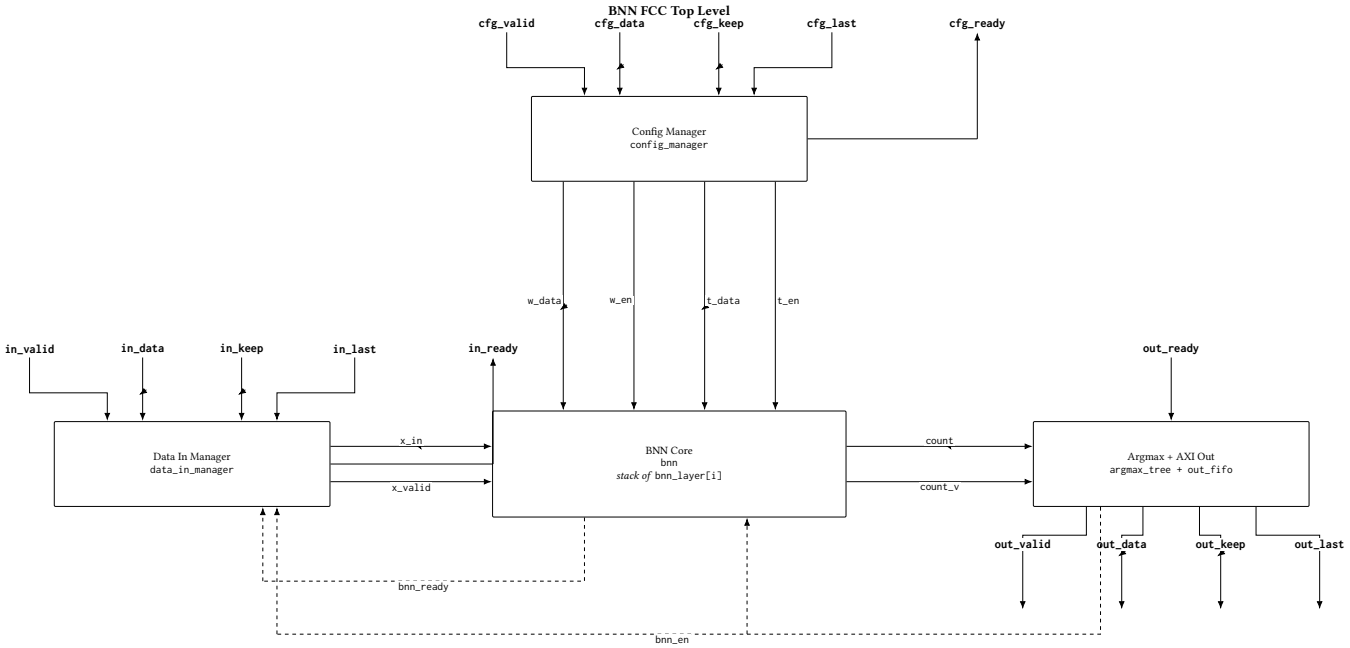

Figure 1: Top-level organization of `rtl/bnn_fcc.sv`. Configuration ingress, activation ingress, the BNN core, and the output side are shown with the main interface signals.

3 Configuration Phase

The configuration phase runs before inference and is responsible for converting the serialized parameter stream into aligned per-layer writes. In this design, that work is intentionally split into two levels. First, the top-level `config_manager` reshapes the incoming byte stream and separates weight payloads from threshold payloads. Then each `bnn_layer` uses a local `config_controller` path to distribute those writes across replicated RAM banks.

The top-level configuration datapath is implemented in `rtl/config_manager.sv` with support from `tkeep_byte_compactor`, `vw_buffer`, `config_manager_parser`, and `config_manager_pad_fsm`. Its job is to accept the external configuration stream, compact away invalid bytes, recover message structure, and emit write-oriented streams that are already separated into the weight and threshold paths.

The front end of this path is byte-oriented because the external interface can arrive with sparse valid-byte patterns. The parser then restores the higher level meaning of the stream by decoding the message header, choosing the weight or threshold destination, and generating the control metadata needed by the padding logic. On the weight side, the pad state machine inserts synthetic bytes so the downstream layer packers can emit aligned write words.

3.1 Parser Control

The parser in `rtl/config_manager_parser.sv` is responsible for recovering message structure from the serialized byte stream. It operates as a small three-state controller: `PARSE_HEADER` consumes the 16-byte header, `PARSE_PAYLOAD` streams the payload bytes, and `PARSE_DONE` waits for the rest of the datapath to drain the current message before the next header begins. In practice, this means the parser owns all protocol-oriented responsibilities: it extracts `msg_type` and `layer_id`, computes the expected payload length from the compile-time layer tables, and asserts `payload_start` exactly when the header has finished.

That separation is intentional. The parser does not try to solve alignment or packer timing at the same time it is decoding the stream. Instead, it emits a clean payload-byte stream plus a small set of control signals and lets the next block deal with write alignment. Splitting those responsibilities into `config_manager_parser` and `config_manager_pad_fsm` made the control easier to implement and easier to reason about, because one FSM is stream/protocol-facing while the other is write-format-facing.

3.2 Padding Control

The padding controller in `rtl/config_manager_pad_fsm.sv` handles the weight-side alignment problem after the parser has already identified the message. Its state machine is likewise explicit: `IDLE` latches the message metadata at `payload_start`, `READ` consumes real payload bytes, `PAD` emits synthetic `0xFF` bytes when a neuron payload does not end on a packer-word boundary, and `DRAIN` flushes the remaining staged bytes before the next message starts. During `PAD`, the FSM asserts `stall` back to the parser so that the protocol side does not advance while synthetic bytes are being inserted.

The `fifo_vr` primitive is what makes that padded-byte path practical. The parser first writes weight payload bytes into a byte-wide `fifo_weights_bytes`, then the pad FSM chooses whether the next byte forwarded downstream is a real FIFO byte or a synthetic padding byte. That single-byte stream is finally repacked by the layer-local `fifo_vr` packer FIFOs into the wider aligned write words expected by each layer. Using the same variable-read FIFO style for both staging and repacking keeps the byte injection mechanism simple: the pad FSM only has to control one downstream byte-write interface, even though the final writes leave as much wider words.

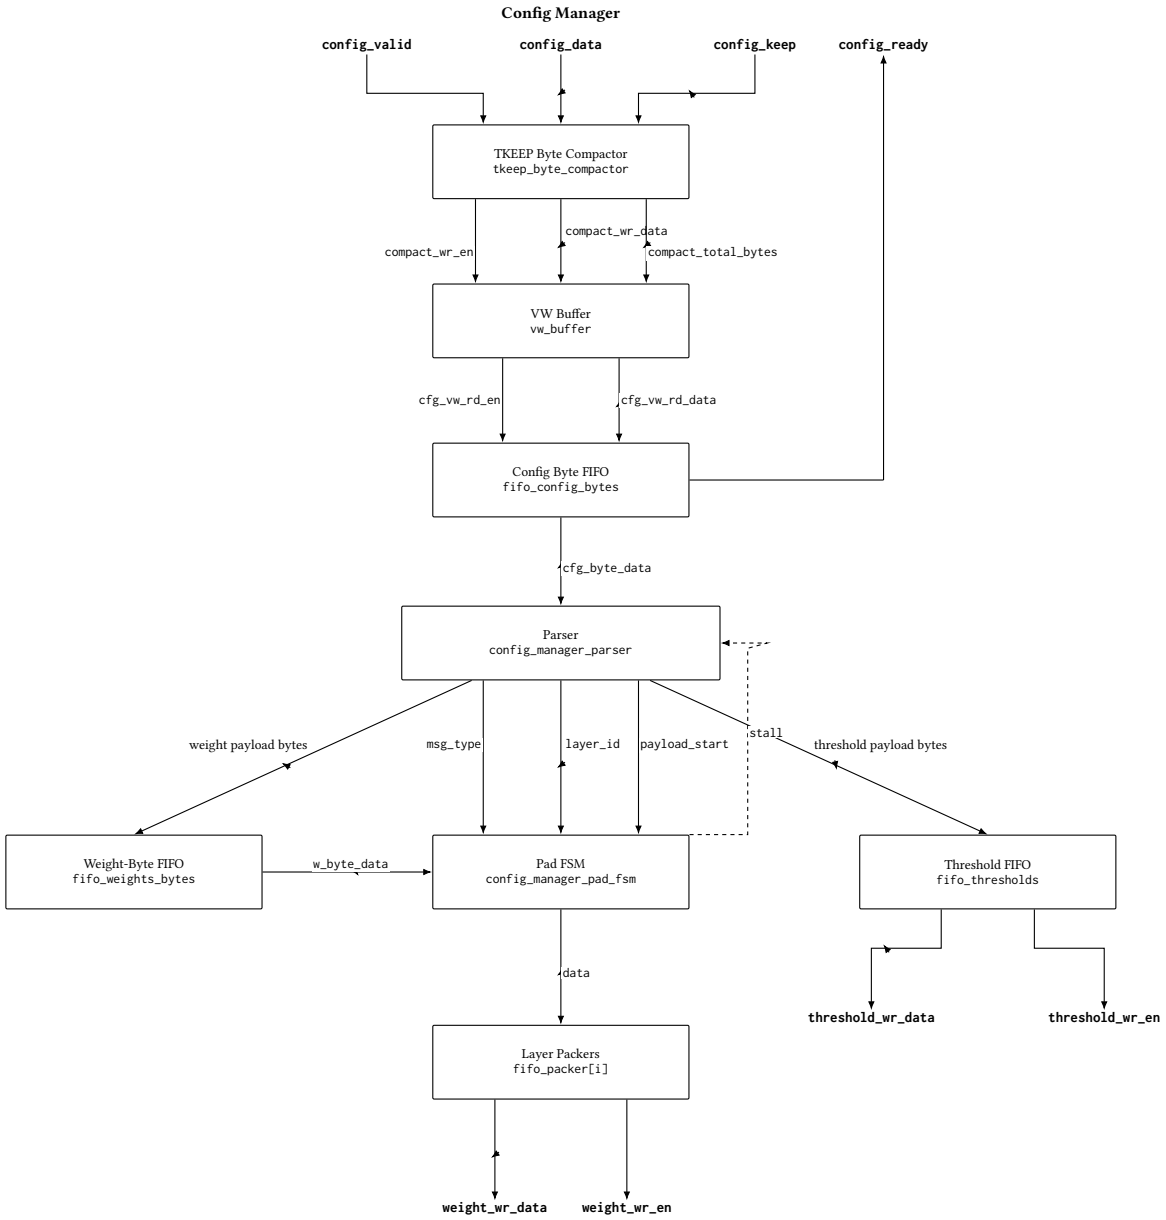

Figure 2: Block diagram of the `config_manager` datapath and its primary interface signals.

3.3 Layer-Local Write Distribution

After the top-level parser finishes that global reshaping step, each `bnn_layer` still has to fan the write stream out to replicated local memories. That is the purpose of `config_controller` and the surrounding staging registers in `rtl/bnn_layer.sv`. The controller generates the bank-select and address signals, while the aligned data stream bypasses the controller datapath and is retimed for fanout.

Those fanout pipeline registers are a central part of the implementation, not a minor detail. They cut the long controller-to-BRAM distribution path into a manageable stage while preserving a steady stream of writes after the pipeline fills. That is why this report treats the local write-distribution block as its own step in the overall configuration phase.

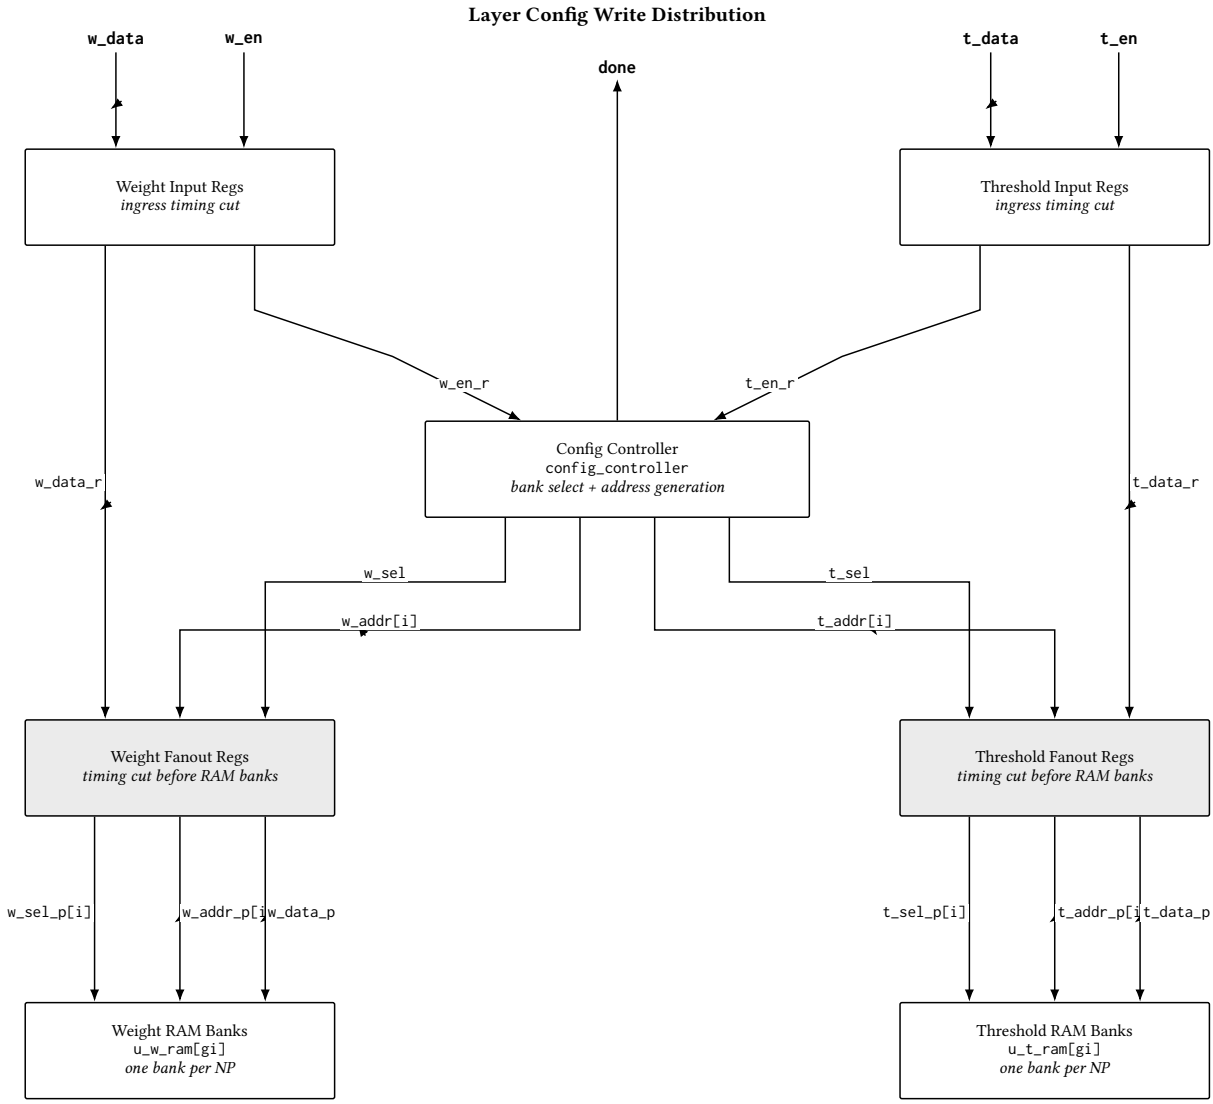

Figure 3: Layer-local config-write distribution around `config_controller`. The light gray fanout pipeline registers mark the timing-cut stage that drives the per-bank weight and threshold RAMs.

4 Compute Phase

Once the model parameters are loaded, the design switches to the compute phase. This phase again has two main pieces. The first is the input-preparation path in `data_in_manager`, which converts the incoming activation stream into the binary words expected by the network. The second is the layer-local forward datapath inside `bnn_layer`, where buffered activations, RAM reads, and neuron-processor outputs are coordinated.

4.1 Input Framing and Packing

`rtl/data_in_manager.sv` mirrors several of the same structural ideas used in the configuration path, but now for activation traffic rather than parameter traffic. The front end again begins with `tkeep_byte_compactor`, because the image stream can arrive with sparse valid-byte patterns. After compaction, local frame-count logic tracks how many real bytes have been accepted in the current image and determines whether the final beat needs zero padding so the frame length becomes a clean multiple of `PARALLEL_INPUTS`. That merged real-data/padding stream is then written through `vw_buffer`, which repacks the byte stream into full internal words.

This block is intentionally separate from the actual layer compute logic. By finishing byte compaction, frame shaping, and word packing at the top level, the `bnn` core sees a regular stream instead of having to carry AXI byte validity and end-of-frame corner cases deep into the runtime datapath.

4.2 Binarization and Staging

After repacking, `data_in_manager` converts each byte into one binary activation bit using the local binarization threshold and writes those bits into the internal `bin_fifo`. This FIFO is more than a convenience buffer: it decouples the image-side repacking path from the BNN-side `bnn_ready`/`bnn_en` handshake, and it provides the backpressure signal that throttles the upstream byte path before the staging queue overflows. In other words, the compute phase still starts with a small stream conditioning problem, and the FIFO is what isolates that problem from the first-layer datapath.

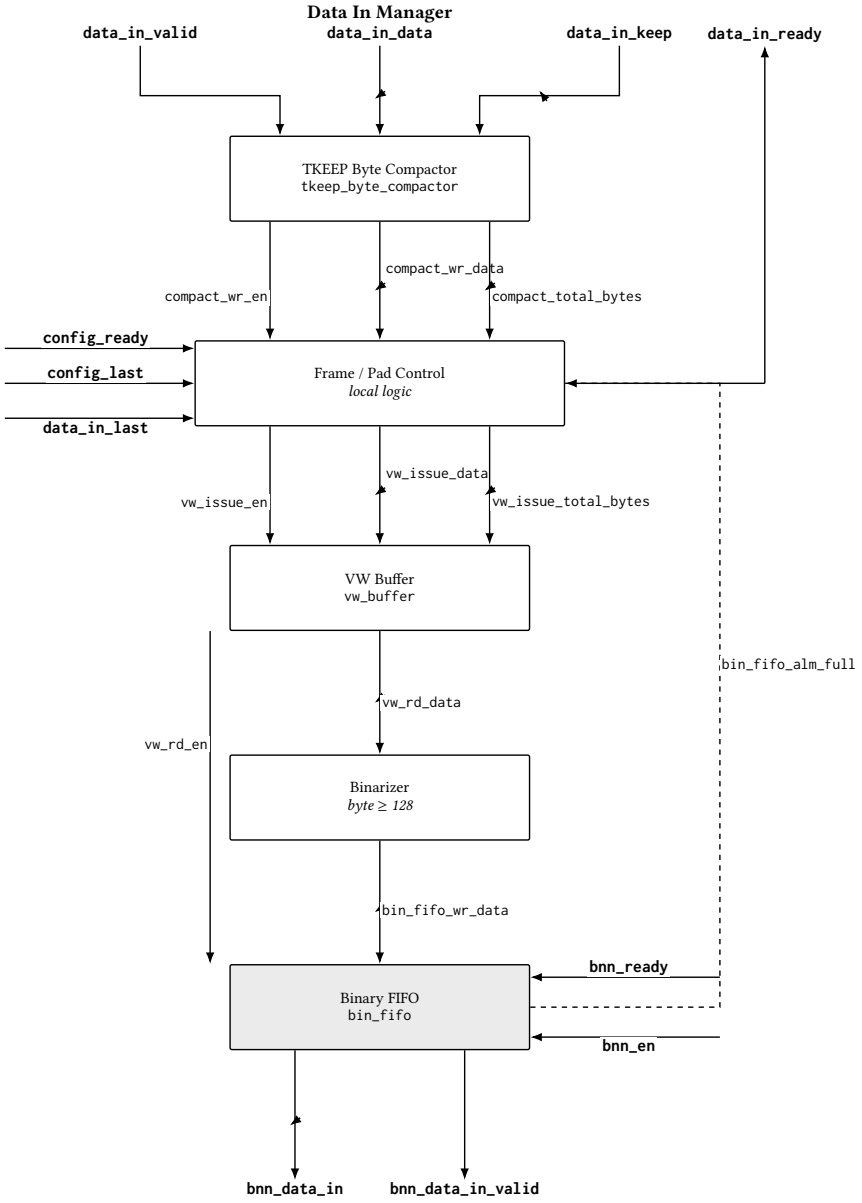

Figure 4: Block diagram of the `data_in_manager` datapath and its primary interface signals. The light gray staging block marks the BNN-facing queue that decouples input reshaping from the core compute flow.

4.3 Replay and Input Alignment

Inside `rtl/bnn_layer.sv`, the forward path begins with the local `replay_buffer`. This block allows one accepted activation frame to be reused across the repeated neuron-group computations required by the layer, instead of demanding the external source resend the same activations for every parallel-neuron slice. The replayed activation word then passes through the explicit alignment delay stage so it reaches the neuron processors in step with the parameter-read pipeline.

That alignment stage matters because the RAM path is not a single-cycle lookup. Between the controller scheduling, local per-bank address registers, and the registered RAM outputs, the weight and threshold operands emerge several cycles after the replay buffer launches the activation word. Keeping the input side explicitly delayed in the diagram makes that timing relationship much easier to understand than burying it inside the controller or RAM descriptions.

4.4 Read Control and Neuron Array

The second half of the compute phase is the parameter-read and evaluation path. `neuron_controller` sequences the shared weight and threshold read addresses, while the banked RAM structures provide one local operand stream per parallel neuron. The gray read-register stages shown in the diagram represent the staging/alignment points that keep the controller side, RAM side, and neuron-processor side synchronized. Finally, the `u_np[gi]` array consumes aligned activations, weights, and thresholds and produces the next layer's output bits along with the accumulator counts used by the last-stage argmax tree.

As in the configuration phase, the implementation is easier to follow because the responsibilities are divided cleanly: replay and alignment handle the activation side, the controller and RAM banks handle operand scheduling, and the neuron processors perform the actual XNOR-accumulate-and-threshold work.

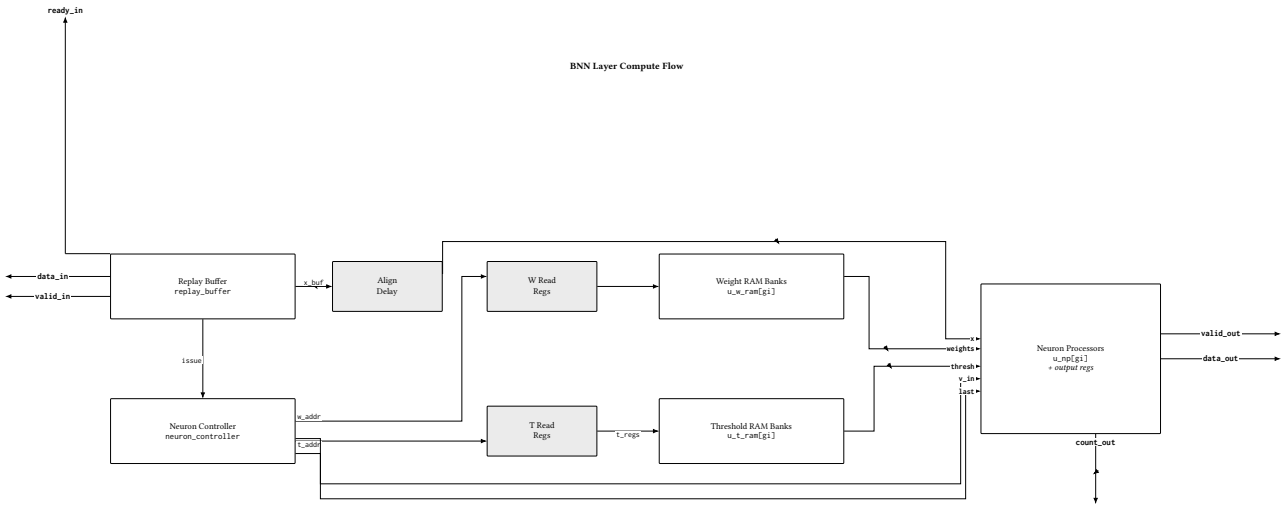

Figure 5: Simplified compute flow inside one `bnn_layer`. The light gray blocks mark the explicit staging and alignment registers in the activation and read-data paths.